student_intervention

April 6, 2016

1 Project 2: Supervised Learning

1.0.1 Building a Student Intervention System

1.1 1. Classification vs Regression

Your goal is to identify students who might need early intervention - which type of supervised machine learning problem is this, classification or regression? Why?

Answer: Classification. This is because the output of our function should tell us weather or not to intervene with a given student. This means constructing a boolean function.

1.2 2. Exploring the Data

Let's go ahead and read in the student dataset first.

*To execute a code cell, click inside it and press **Shift+Enter**.*

```
In [1]: # Import libraries
        import numpy as np
        import pandas as pd

In [2]: # Read student data
        student_data = pd.read_csv("student-data.csv")
        print "Student data read successfully!"
        # Note: The last column 'passed' is the target/label, all other are feature columns
```

```
Student data read successfully!
```

Now, can you find out the following facts about the dataset? - Total number of students - Number of students who passed - Number of students who failed - Graduation rate of the class (%) - Number of features

*Use the code block below to compute these values. Instructions/steps are marked using **TODOs**.*

```
In [3]: # TODO: Compute desired values - replace each '?' with an appropriate expression/function call
        n_students = student_data.index.size
        n_features = student_data.columns.size - 1
        n_passed = student_data[ student_data['passed'] == 'yes' ].index.size
        n_failed = n_students - n_passed
        grad_rate = n_passed / float(n_students)
        print "Total number of students: {}".format(n_students)
        print "Number of students who passed: {}".format(n_passed)
        print "Number of students who failed: {}".format(n_failed)
        print "Number of features: {}".format(n_features)
        print "Graduation rate of the class: {:.2f}%".format(grad_rate)
```

```
Total number of students: 395
Number of students who passed: 265
Number of students who failed: 130
Number of features: 30
Graduation rate of the class: 0.67%
```

1.3 3. Preparing the Data

In this section, we will prepare the data for modeling, training and testing.

1.3.1 Identify feature and target columns

It is often the case that the data you obtain contains non-numeric features. This can be a problem, as most machine learning algorithms expect numeric data to perform computations with.

Let's first separate our data into feature and target columns, and see if any features are non-numeric. **Note**: For this dataset, the last column (`passed`) is the target or label we are trying to predict.

```
In [4]: # Extract feature (X) and target (y) columns
        feature_cols = list(student_data.columns[:-1])  # all columns but last are features
        target_col = student_data.columns[-1]  # last column is the target/label
        print "Feature column(s):-\n{}".format(feature_cols)
        print "Target column: {}".format(target_col)

        X_all = student_data[feature_cols]  # feature values for all students
        y_all = student_data[target_col]  # corresponding targets/labels
        print "\nFeature values:-"
        print X_all.head()  # print the first 5 rows
```

```
Feature column(s):-
['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'Mjob', 'Fjob', 'reason', 'gua
Target column: passed
```

```
Feature values:-
  school sex  age address famsize Pstatus  Medu  Fedu     Mjob      Fjob  \
0     GP   F   18       U     GT3       A     4     4  at_home   teacher
1     GP   F   17       U     GT3       T     1     1  at_home     other
2     GP   F   15       U     LE3       T     1     1  at_home     other
3     GP   F   15       U     GT3       T     4     2   health  services
4     GP   F   16       U     GT3       T     3     3    other     other

       ...   higher internet  romantic  famrel  freetime goout Dalc Walc health  \
0      ...      yes       no        no       4         3     4    1    1      3
1      ...      yes      yes        no       5         3     3    1    1      3
2      ...      yes      yes        no       4         3     2    2    3      3
3      ...      yes      yes       yes       3         2     2    1    1      5
4      ...      yes       no        no       4         3     2    1    2      5

   absences
0         6
1         4
2        10
3         2
4         4

[5 rows x 30 columns]
```

1.3.2 Preprocess feature columns

As you can see, there are several non-numeric columns that need to be converted! Many of them are simply yes/no, e.g. `internet`. These can be reasonably converted into 1/0 (binary) values.

Other columns, like `Mjob` and `Fjob`, have more than two values, and are known as *categorical variables*. The recommended way to handle such a column is to create as many columns as possible values (e.g. `Fjob_teacher`, `Fjob_other`, `Fjob_services`, etc.), and assign a `1` to one of them and `0` to all others.

These generated columns are sometimes called *dummy variables*, and we will use the `pandas.get_dummies()` function to perform this transformation.

```python
In [5]: # Preprocess feature columns
        def preprocess_features(X):
            outX = pd.DataFrame(index=X.index)  # output dataframe, initially empty

            # Check each column
            for col, col_data in X.iteritems():
                # If data type is non-numeric, try to replace all yes/no values with 1/0
                if col_data.dtype == object:
                    col_data = col_data.replace(['yes', 'no'], [1, 0])
                # Note: This should change the data type for yes/no columns to int

                # If still non-numeric, convert to one or more dummy variables
                if col_data.dtype == object:
                    col_data = pd.get_dummies(col_data, prefix=col)  # e.g. 'school' => 'school_GP', 's

                outX = outX.join(col_data)  # collect column(s) in output dataframe

            return outX

        X_all = preprocess_features(X_all)
        y_all = y_all.replace(['yes','no'],[1,0])
        print "Processed feature columns ({}):-\n{}".format(len(X_all.columns), list(X_all.columns))
```

```
Processed feature columns (48):-
['school_GP', 'school_MS', 'sex_F', 'sex_M', 'age', 'address_R', 'address_U', 'famsize_GT3', 'famsize_LE3'
```

1.3.3 Split data into training and test sets

So far, we have converted all *categorical* features into numeric values. In this next step, we split the data (both features and corresponding labels) into training and test sets.

```python
In [6]: # First, decide how many training vs test samples you want
        num_all = student_data.shape[0]  # same as len(student_data)
        num_train = 300  # about 75% of the data
        num_test = num_all - num_train

        # TODO: Then, select features (X) and corresponding labels (y) for the training and test sets
        # Note: Shuffle the data or randomly select samples to avoid any bias due to ordering in the da
        from sklearn.cross_validation import train_test_split
        X_train, X_test, y_train, y_test = train_test_split( X_all , y_all , test_size = num_test, rand

        print "Training set: {} samples".format(X_train.shape[0])
        print "Test set: {} samples".format(X_test.shape[0])
        # Note: If you need a validation set, extract it from within training data
```

```
Training set: 300 samples
Test set: 95 samples
```

1.4   4. Training and Evaluating Models

Choose 3 supervised learning models that are available in scikit-learn, and appropriate for this problem. For each model:

- What is the theoretical O(n) time & space complexity in terms of input size?
- What are the general applications of this model? What are its strengths and weaknesses?
- Given what you know about the data so far, why did you choose this model to apply?
- Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.

Produce a table showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.
Note: You need to produce 3 such tables - one for each model.

1.5   Answer

In the following n_f denotes the number of featurs, and n_s denotes the number of samples.

1.5.1   Decision Trees :

- **Complexity:** Training time is $O(n_f n_s^2 \log(n_s))$ in time. Query time is $O(\log(n_f))$. The storage cost of the tree is $O(n_f^2)$.
- **Applications:** Decision trees are useful when individual features can efficiently split the data. For example, in the case of determining if it will rain, knowing weather or not there are clouds in the sky is a feature worth splitting your algorithm on.
- **Strengths:** Decisision trees are human interpretable (assuming the features are), and they are space efficient.
- **Weaknesses:** The training algorithms are always suboptimal because the building of an optimal decision tree is an non-polynomial time problem (as far as we know). One side-effect of this sub-optimality is that the classifier is unstable with respect to perturbations in the data.
- **Why should we condiser it?:** We should consider decision trees for the strengths listed above, and the fact that much of our data is binary, and thus well suited to a Decision tree.

Performance:

training set size	training time	prediction times (training/test)	F1 scores (training/test)
100	0.002	0.000 / 0.000	1.000 / 0.626
200	0.003	0.000 / 0.001	1.000 / 0.753
300	0.004	0.001 / 0.001	1.000 / 0.672

1.5.2   Naive Bayes :

- **Complexity:** Training time is $O(1)$ in time and $O(n)$ in space. Querying is (worst case scenario) $O(n)$ in time, and $O(1)$ in additional space.
- **Applications:** Assumes that each piece of data is obtained independently of the others (in the sense of probability theory). If this is in fact the case, then this algorithm might even be optimal if one can choose the correct model for $P(y \mid x)$. Verifying independence is not possible, but you can certainly

accumulate evidence for this assumption.
- **Strengths:** Fast training time. If you are only going to make a single query this might be a very efficient algorithm.
- **Weaknesses:** If your samples are not independent, than this algorithm will not be good at determining the underlying (and possibly lower dimensional) structure of the data.
- **Why did you choose this?** I have not tested for independence, however if the samples are independent, than this algorithm might beat the other algorithms I've chosen to test.

Performance:

training set size	training time	prediction times (training/test)	F1 scores (training/test)
100	0.002	0.001 / 0.000	0.804 / 0.756
200	0.002	0.001 / 0.000	0.804 / 0.756
300	0.002	0.010 / 0.007	0.804 / 0.756

1.5.3 K-nearest neighbors :

- **Complexity:** Training is $O(1)$ in time and $O(n)$ in space. Querying is $O(log(n) + k)$ in time if the feature vectors are sorted and $O(n + k)$ otherwise. Querying is $O(1)$ additional space, in time.
- **Appplictions:** Such an algorithm is biased towards smoother data, and useful when the numer of queries is small.
- **Strengths:** Very fast training time. Highly flexible, as one can choose the distance function as well as k.
- **Weakness:** Perhaps it is too flexible. It takes a lot of work to optimize the parameters.
- **Reason for choosing this algorithm:** Unlike Naive Bayes, no such independence assumption is made. Moreover, the metric can be chosen to specialize the algorithm to data sets with Boolean features. The flexibility of this algorithms gives it more knobs to tweak and it might have superior performance to the other algorithms after doing a grid search over the parameter space.

Performance:

training set size	training time	prediction times (training/test)	F1 scores (training/test)
100	0.001	0.030 / 0.011	0.824 / 0.757
200	0.001	0.033 / 0.008	0.824 / 0.757
300	0.001	0.024 / 0.009	0.824 / 0.757

```
In [7]: # Train a model
        import time

        def train_classifier(clf, X_train, y_train):
            print "Training {}...".format(clf.__class__.__name__)
```

```python
        start = time.time()
        clf.fit(X_train, y_train)
        end = time.time()
        print "Done!\nTraining time (secs): {:.3f}".format(end - start)

In [8]: # Predict on training set and compute F1 score
        from sklearn.metrics import f1_score

        def predict_labels(clf, features, target):
            print "Predicting labels using {}...".format(clf.__class__.__name__)
            start = time.time()
            y_pred = clf.predict(features)
            end = time.time()
            print "Done!\nPrediction time (secs): {:.3f}".format(end - start)
            return f1_score(target.values, y_pred, pos_label=1)

In [9]: # Train and predict using different training set sizes
        def train_predict(clf, X_train, y_train, X_test, y_test):
            print "------------------------------------------"
            print "Training set size: {}".format(len(X_train))
            train_classifier(clf, X_train, y_train)
            print "F1 score for training set: {}".format(predict_labels(clf, X_train, y_train))
            print "F1 score for test set: {}".format(predict_labels(clf, X_test, y_test))

        # TODO: Run the helper function above for desired subsets of training data
        # Note: Keep the test set constant

In [10]: # TODO: Train and predict using two other models
         # TODO: Choose a model, import it and instantiate an object
         num_train = 200

         from sklearn.tree import DecisionTreeClassifier
         dt_clf = DecisionTreeClassifier()
         f1_dt = train_predict( dt_clf , X_train[0:num_train], y_train[0:num_train], X_test, y_test )
         from sklearn.naive_bayes import BernoulliNB
         nb_clf = BernoulliNB()
         f1_nb = train_predict(nb_clf , X_train, y_train, X_test, y_test)
         from sklearn.neighbors import KNeighborsClassifier
         knn_clf = KNeighborsClassifier(metric='jaccard')
         f1_knn = train_predict( knn_clf , X_train, y_train , X_test, y_test)

------------------------------------------
Training set size: 200
Training DecisionTreeClassifier...
Done!
Training time (secs): 0.004
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.000
F1 score for training set: 1.0
Predicting labels using DecisionTreeClassifier...
Done!
Prediction time (secs): 0.000
F1 score for test set: 0.732824427481
------------------------------------------
```

```
Training set size: 300
Training BernoulliNB...
Done!
Training time (secs): 0.013
Predicting labels using BernoulliNB...
Done!
Prediction time (secs): 0.002
F1 score for training set: 0.803695150115
Predicting labels using BernoulliNB...
Done!
Prediction time (secs): 0.000
F1 score for test set: 0.755555555556
------------------------------------------
Training set size: 300
Training KNeighborsClassifier...
Done!
Training time (secs): 0.005
Predicting labels using KNeighborsClassifier...
Done!
Prediction time (secs): 0.025
F1 score for training set: 0.824295010846
Predicting labels using KNeighborsClassifier...
Done!
Prediction time (secs): 0.008
F1 score for test set: 0.757142857143
```

1.6 5. Choosing the Best Model

- Based on the experiments you performed earlier, in 1-2 paragraphs explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?
- In 1-2 paragraphs explain to the board of supervisors in layman's terms how the final model chosen is supposed to work (for example if you chose a Decision Tree or Support Vector Machine, how does it make a prediction).
- Fine-tune the model. Use Gridsearch with at least one important parameter tuned and with at least 3 settings. Use the entire training set for this.
- What is the model's final F1 score?

Choice of model: I've come to the conclusion that the k-nearest neighbors algorithm is best suited to the data. The F1 score on the test set of 0.797 is the highest among the other F1 scores which were tested (0.789 for generic support vector machines and 0.687 for linear support vector machines). A high F1 score indicates: 1. When a student is likely to fail, our algorithm will suggest intervention (i.e. high recall) 2. When the algorithm suggests intervention, it is usually the case that he or she is in danger of failing (i.e. high precision)

These are both things which we'd like to maximize in our intervention algorithm, and the k nearest neighbors algorithm achieved the highest score among the explored alternatives. Moreover, the gap between the F1 scores on the training and testing sets is least for the k-nearest neighbors algorithm, which suggests that it uses the given data in a generalizable way (i.e. it is least guilty of overfitting). A potential draw-back of the k-nearest neighbors algorithm is that when there are many many records (e.g. 10,000) then the time it takes for the computer to determine intervention might get large (i.e. prediction for k-nearest neighbors does not scale well with data size). This is not the case with support vector machines, which were the alternative algorithms which I tested. However, the current data set consists of 395 students, a relatively small number. Given 400 students every year, this will become a problem in about 25 years. By 2041 I expect that our concerns will shift towards hunting replicants who've escaped the off-world colonies. It seems misplaced to worry about the scaling of this algorithm at this time.

How the model works: The k-nearest neighbors algorithm works as follows. Given a student, whom we'd like to determine weather or not to intervene, the algorithm will find k students in our data who are most similar (where similar is defined by a customizable distance function). The algorithm will use the outcomes of these nearby students to compute a probability of intervention for the student in question, perhaps by averaging. This simple interpretation is another advantage of the k-nearest neighbors algorithm over support vector machines.

** Fine tuning:** One of the free parameters is what we use to measure the similarity between students. This is given by a distance function. We can also vary k, the number of neighbors. Below we consider 6 possible metrics, and k = 1,2,4,8,16,32,64. It appears the optimal metric (in terms of F1 scores) is the "kulinski" distance with and optimal k of 32. The F1 score of the optimal classifier is 0.812.

```
In [11]: # TODO: Fine-tune your model and report the best F1 score, use GridSearch
         from sklearn.grid_search import GridSearchCV
         metrics = ['jaccard','matching','dice','kulsinski','rogerstanimoto','russellrao']
         k_list = map( lambda p: 2**p , range(6))
         parameters = { 'n_neighbors':k_list, 'weights':['uniform','distance'], 'metric':metrics}
         model = KNeighborsClassifier()
         from sklearn.metrics import make_scorer
         scorer = make_scorer( f1_score )
         clf = GridSearchCV( model , parameters , scoring = "f1" )
         clf.fit( X_all , y_all )
         print clf.best_estimator_
         print 'f1 score = {:f}'.format(f1_score( clf.predict(X_all) , y_all))

KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='kulsinski',
          metric_params=None, n_neighbors=32, p=2, weights='uniform')
f1 score = 0.811820
```