

student_intervention

April 7, 2016

1 Project 2: Supervised Learning

1.0.1 Building a Student Intervention System

1.1 1. Classification vs Regression

Your goal is to identify students who might need early intervention - which type of supervised machine learning problem is this, classification or regression? Why?

Answer: Classification. This is because the output of our function should tell us weather or not to intervene with a given student. This means constructing a boolean function.

1.2 2. Exploring the Data

Let's go ahead and read in the student dataset first.

*To execute a code cell, click inside it and press **Shift+Enter**.*

```
In [2]: # Import libraries
import numpy as np
import pandas as pd
```

```
In [3]: # Read student data
student_data = pd.read_csv("student-data.csv")
print "Student data read successfully!"
# Note: The last column 'passed' is the target/label, all other are feature columns
```

Student data read successfully!

Now, can you find out the following facts about the dataset? - Total number of students - Number of students who passed - Number of students who failed - Graduation rate of the class (%) - Number of features

*Use the code block below to compute these values. Instructions/steps are marked using **TODOs**.*

```
In [4]: # TODO: Compute desired values - replace each '?' with an appropriate expression/function call
n_students = student_data.index.size
n_features = student_data.columns.size - 1
n_passed = student_data[ student_data['passed'] == 'yes' ].index.size
n_failed = n_students - n_passed
grad_rate = 100*float(n_passed) / float(n_students)
print "Total number of students: {}".format(n_students)
print "Number of students who passed: {}".format(n_passed)
print "Number of students who failed: {}".format(n_failed)
print "Number of features: {}".format(n_features)
print "Graduation rate of the class: {:.2f}%".format(grad_rate)
```

Total number of students: 395

Number of students who passed: 265

Number of students who failed: 130

Number of features: 30

Graduation rate of the class: 67.09%

1.3 3. Preparing the Data

In this section, we will prepare the data for modeling, training and testing.

1.3.1 Identify feature and target columns

It is often the case that the data you obtain contains non-numeric features. This can be a problem, as most machine learning algorithms expect numeric data to perform computations with.

Let's first separate our data into feature and target columns, and see if any features are non-numeric.

Note: For this dataset, the last column ("passed") is the target or label we are trying to predict.

```
In [5]: # Extract feature (X) and target (y) columns
feature_cols = list(student_data.columns[:-1]) # all columns but last are features
target_col = student_data.columns[-1] # last column is the target/label
print "Feature column(s):-\n{}".format(feature_cols)
print "Target column: {}".format(target_col)

X_all = student_data[feature_cols] # feature values for all students
y_all = student_data[target_col] # corresponding targets/labels
print "\nFeature values:-"
print X_all.head() # print the first 5 rows
```

Feature column(s):-

['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'Mjob', 'Fjob', 'reason', 'guardian']

Target column: passed

Feature values:-

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	
1	GP	F	17	U	GT3	T	1	1	at_home	other	
2	GP	F	15	U	LE3	T	1	1	at_home	other	
3	GP	F	15	U	GT3	T	4	2	health	services	
4	GP	F	16	U	GT3	T	3	3	other	other	

	...	higher	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	\
0	...	yes	no	no	4	3	4	1	1	3	
1	...	yes	yes	no	5	3	3	1	1	3	
2	...	yes	yes	no	4	3	2	2	3	3	
3	...	yes	yes	yes	3	2	2	1	1	5	
4	...	yes	no	no	4	3	2	1	2	5	

absences

0	6
1	4
2	10
3	2
4	4

[5 rows x 30 columns]

1.3.2 Preprocess feature columns

As you can see, there are several non-numeric columns that need to be converted! Many of them are simply yes/no, e.g. `internet`. These can be reasonably converted into 1/0 (binary) values.

Other columns, like Mjob and Fjob, have more than two values, and are known as *categorical variables*. The recommended way to handle such a column is to create as many columns as possible values (e.g. Fjob_teacher, Fjob_other, Fjob_services, etc.), and assign a 1 to one of them and 0 to all others.

These generated columns are sometimes called *dummy variables*, and we will use the `pandas.get_dummies()` function to perform this transformation.

```
In [6]: # Preprocess feature columns
def preprocess_features(X):
    outX = pd.DataFrame(index=X.index) # output dataframe, initially empty

    # Check each column
    for col, col_data in X.iteritems():
        # If data type is non-numeric, try to replace all yes/no values with 1/0
        if col_data.dtype == object:
            col_data = col_data.replace(['yes', 'no'], [1, 0])
            # Note: This should change the data type for yes/no columns to int

        # If still non-numeric, convert to one or more dummy variables
        if col_data.dtype == object:
            col_data = pd.get_dummies(col_data, prefix=col) # e.g. 'school' => 'school_GP', 'school_MS'

    outX = outX.join(col_data) # collect column(s) in output dataframe

    return outX

X_all = preprocess_features(X_all)
y_all = y_all.replace(['yes', 'no'], [1, 0])
print "Processed feature columns ({}):-\n{}".format(len(X_all.columns), list(X_all.columns))
```

Processed feature columns (48):-

['school_GP', 'school_MS', 'sex_F', 'sex_M', 'age', 'address_R', 'address_U', 'famsize_GT3', 'famsize_LE3']

1.3.3 Split data into training and test sets

So far, we have converted all *categorical* features into numeric values. In this next step, we split the data (both features and corresponding labels) into training and test sets.

```
In [7]: # First, decide how many training vs test samples you want
num_all = student_data.shape[0] # same as len(student_data)
num_train = 300 # about 75% of the data
num_test = num_all - num_train

# TODO: Then, select features (X) and corresponding labels (y) for the training and test sets
# Note: Shuffle the data or randomly select samples to avoid any bias due to ordering in the data
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X_all , y_all , test_size = num_test, random_state=0)

print "Training set: {} samples".format(X_train.shape[0])
print "Test set: {} samples".format(X_test.shape[0])
# Note: If you need a validation set, extract it from within training data
```

Training set: 300 samples

Test set: 95 samples

1.4 4. Training and Evaluating Models

Choose 3 supervised learning models that are available in scikit-learn, and appropriate for this problem. For each model:

- What is the theoretical $O(n)$ time & space complexity in terms of input size?
- What are the general applications of this model? What are its strengths and weaknesses?
- Given what you know about the data so far, why did you choose this model to apply?
- Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.

Produce a table showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.

Note: You need to produce 3 such tables - one for each model.

1.5 Answer

In the following n_f denotes the number of features, and n denotes the number of samples.

1.5.1 Decision Trees :

- **Complexity:** I'm not sure how to compute the complexities for a generic decision tree. However, if you assume that the data can be split perfectly at each level of the tree then you can calculate the following. Training time is $O(n_f n_s^2 \log(n_s))$ in time. Query time is $O(\log(n_f))$. The storage cost of the tree is $O(n_f^2)$.
- **Applications:** Decision trees are useful when individual features can efficiently split the data. For example, in the case of determining if it will rain, knowing weather or not there are clouds in the sky is a feature worth splitting your algorithm on.
- **Strengths:** Decision trees are human interpretable (assuming the features are), and they are space efficient. Moreover, they have very low bias. In the language of function spaces, the decision trees are generated by rectangular indicator functions on the feature space. The function space generated by these box-functions is just the space of measurable function on the feature space (with range in the target set). In otherwords, Decision trees have lots of expressive power.
- **Weaknesses:** The training algorithms are always suboptimal because the building of an optimal decision tree is a non-polynomial time problem (as far as we know). One side-effect of this sub-optimality is that the classifier is unstable with respect to perturbations in the data. Moreover, Decision trees are subject to overfitting. This is especially the case when the leaves contain only one member of the training set, which is the natural place where training ends unless you actively prevent this outcome through pruning or other variance reduction measures.
- **Why should we consider it?:** We should consider decision trees for the strengths listed above, and the fact that much of our data is binary, and thus well suited to the decision tree paradigm (if in fact a good decision tree can be found).

Performance:

training set size	training time	prediction times (training/test)	F1 scores
			(train- ing/test)
100	0.002	0.001 / 0.000	1.000 / 0.603
200	0.003	0.000 / 0.000	1.000 / 0.716
300	0.004	0.000 / 0.000	1.000 / 0.600

1.5.2 Naive Bayes :

- **Complexity:** Training time is $O(1)$ in time and $O(n)$ in space. Querying is (worst case scenario) $O(n)$ in time, and $O(1)$ in additional space.
- **Applications:** Naive Bayes assumes that each piece of data is obtained independently of the others (in the sense of probability theory). If this is in fact the case, then this algorithm might even be optimal if one can choose the correct model for $P(y | x)$. Verifying independence is typically not possible, but you can certainly accumulate evidence for this assumption.
- **Strengths:** Fast training time. If you are only going to make a single query this might be a very efficient algorithm. Additionally, if the independence assumption is true, Naive Bayes can have very low bias.
- **Weaknesses:** If your samples are not independent, then this algorithm will not be good at determining the underlying (and possibly lower dimensional) structure of the data. In other words, when the independence assumption is false, Naive Bayes will yield a predictor with high bias. Moreover, there is potential for overfitting, as Naive Bayes is a lazy learner and thus prone to simply memorizing the training set rather than generalizing it.
- **Why did you choose this?** I have not tested for independence, however if the samples are independent, then this algorithm might beat the other algorithms I've chosen to test. Many educators implicitly hold this independence assumption (at least in some high level sense) when they speak of each student being judged independently, so perhaps there is some merit to it.

Performance:

training set size	training time	prediction times (training/test)	F1 scores
			(train- ing/test)
100	0.009	0.000 / 0.000	0.829 / 0.794
200	0.001	0.001 / 0.000	0.827 / 0.779
300	0.009	0.010 / 0.007	0.804 / 0.756

1.5.3 K-nearest neighbors :

- **Complexity:** Training is $O(1)$ in time and $O(n)$ in space. Querying is $O(\log(n) + k)$ in time if the feature vectors are sorted and $O(n + k)$ otherwise. Querying is $O(1)$ additional space, in time.
- **Applications:** Such an algorithm is biased towards smoother data, and useful when the number of queries is small.
- **Strengths:** Very fast training time. Highly flexible, as one can choose the distance function as well as k .
- **Weakness:** Perhaps it is too flexible. It takes a lot of work to optimize the parameters. When k is small, then the training set might be too influential, and we might overfit (i.e. high variance). When k is large we might suffer from high bias.
- **Reason for choosing this algorithm:** Unlike Naive Bayes, no such independence assumption is made. Moreover, the metric can be chosen to specialize the algorithm to data sets with Boolean features. The flexibility of this algorithm gives it more knobs to tweak and it might have superior performance to the other algorithms after doing a grid search over the parameter space.

Performance:

training set size	training time	prediction times (training/test)	F1 scores (train- ing/test)
100	0.001	0.003 / 0.003	0.766 / 0.730
200	0.001	0.014 / 0.006	0.841 / 0.727
300	0.006	0.028 / 0.012	0.824 / 0.757

```
In [8]: # Train a model
import time

def train_classifier(clf, X_train, y_train):
    print "Training {}...".format(clf.__class__.__name__)
    start = time.time()
    clf.fit(X_train, y_train)
    end = time.time()
    print "Training time (secs): {:.3f}".format(end - start)

In [9]: # Predict on training set and compute F1 score
from sklearn.metrics import f1_score

def predict_labels(clf, features, target):
    #print "Predicting labels using {}...".format(clf.__class__.__name__)
    start = time.time()
    y_pred = clf.predict(features)
    end = time.time()
    print "Prediction time (secs): {:.3f}".format(end - start)
    return f1_score(target.values, y_pred, pos_label=1)

In [21]: # Train and predict using different training set sizes
def train_predict(clf, X_train, y_train, X_test, y_test):
    print "-----"
    print "Training set size: {}".format(len(X_train))
    train_classifier(clf, X_train, y_train)
    print "F1 score for training set: {}".format(predict_labels(clf, X_train, y_train))
    print "F1 score for test set: {}".format(predict_labels(clf, X_test, y_test))

    # TODO: Run the helper function above for desired subsets of training data
    # Note: Keep the test set constant

In [22]: # TODO: Train and predict using two other models
# TODO: Choose a model, import it and instantiate an object
num_train_arr = [100,200,300]

from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn.neighbors import KNeighborsClassifier

dt_clf = DecisionTreeClassifier(random_state = 32)
nb_clf = BernoulliNB()
knn_clf = KNeighborsClassifier(metric='jaccard')
```

```

        clf_arr = [dt_clf,nb_clf,knn_clf]

        for clf in clf_arr:
            for num_train in num_train_arr:
                f1_dt = train_predict( clf , X_train[0:num_train], y_train[0:num_train], X_test, y_test)

-----
Training set size: 100
Training DecisionTreeClassifier...
Training time (secs): 0.002
Prediction time (secs): 0.001
F1 score for training set: 1.0
Prediction time (secs): 0.000
F1 score for test set: 0.603448275862
-----
Training set size: 200
Training DecisionTreeClassifier...
Training time (secs): 0.003
Prediction time (secs): 0.000
F1 score for training set: 1.0
Prediction time (secs): 0.000
F1 score for test set: 0.716417910448
-----
Training set size: 300
Training DecisionTreeClassifier...
Training time (secs): 0.004
Prediction time (secs): 0.000
F1 score for training set: 1.0
Prediction time (secs): 0.000
F1 score for test set: 0.6
-----
Training set size: 100
Training BernoulliNB...
Training time (secs): 0.009
Prediction time (secs): 0.000
F1 score for training set: 0.828571428571
Prediction time (secs): 0.000
F1 score for test set: 0.794326241135
-----
Training set size: 200
Training BernoulliNB...
Training time (secs): 0.001
Prediction time (secs): 0.001
F1 score for training set: 0.826666666667
Prediction time (secs): 0.000
F1 score for test set: 0.779411764706
-----
Training set size: 300
Training BernoulliNB...
Training time (secs): 0.009
Prediction time (secs): 0.001
F1 score for training set: 0.803695150115
Prediction time (secs): 0.001
F1 score for test set: 0.755555555556

```

```

-----
Training set size: 100
Training KNeighborsClassifier...
Training time (secs): 0.001
Prediction time (secs): 0.003
F1 score for training set: 0.765957446809
Prediction time (secs): 0.003
F1 score for test set: 0.729927007299
-----

Training set size: 200
Training KNeighborsClassifier...
Training time (secs): 0.001
Prediction time (secs): 0.013
F1 score for training set: 0.84076433121
Prediction time (secs): 0.006
F1 score for test set: 0.727272727273
-----

Training set size: 300
Training KNeighborsClassifier...
Training time (secs): 0.006
Prediction time (secs): 0.025
F1 score for training set: 0.824295010846
Prediction time (secs): 0.008
F1 score for test set: 0.757142857143
-----

```

1.6 5. Choosing the Best Model

- Based on the experiments you performed earlier, in 1-2 paragraphs explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?
- In 1-2 paragraphs explain to the board of supervisors in layman's terms how the final model chosen is supposed to work (for example if you chose a Decision Tree or Support Vector Machine, how does it make a prediction).
- Fine-tune the model. Use Gridsearch with at least one important parameter tuned and with at least 3 settings. Use the entire training set for this.
- What is the model's final F1 score?

Choice of model: I've come to the conclusion that the k-nearest neighbors algorithm is best suited to the data. The F_1 -scores for Naive Bayes and Decision trees seem to indicate both algorithms are suffering over-fit when we use the out-of-the-box implementations. This is because the F_1 -scores seem to vary by more than 0.15 when comparing the score of the test set vs the training set. However, in the case of a training set of size 100, the F_1 -score of k-nearest neighbors only varies by .036, which suggests that we are not suffering from over-fit. Moreover, the F_1 score on the test set is as high for k-nearest neighbors as it is for the other algorithms tested. A high F_1 -score indicates: 1. When a student is likely to fail, our algorithm will suggest intervention (i.e. high recall) 2. When the algorithm suggests intervention, it is usually the case that he or she is in danger of failing (i.e. high precision)

These are both things which we'd like to maximize in our intervention algorithm.

A potential draw-back of the k-nearest neighbors algorithm is that when there are many many records (e.g. 10,000) then the time it takes for the computer to determine intervention might get large (i.e. prediction for k-nearest neighbors does not scale well with data size). This is not the case with support vector machines, which were the alternative algorithms which I tested. However, the current data set consists of 395 students, a relatively small number. Given 400 students every year, this will become a problem in about 25 years. Actually, this is a decent example of an embarassingly parallelizable algorithm. So this concern is really mute if one is willing to take advantage of parallelization.

How the model works: The k-nearest neighbors algorithm works as follows. Given a student, whom we'd like to determine weather or not to intervene, the algorithm will find k students in our data who are most similar (where similar is defined by a customizable distance function). The algorithm will use the outcomes of these nearby students to compute a probability of intervention for the student in question, perhaps by a simple vote.

**** Fine tuning:**** One of the free parameters is what we use to measure the similarity between students. This is given by a distance function. We can also vary k, the number of neighbors. Below we consider 6 possible metrics, and k = 1,2,4,8,16,32,64. It appears the optimal metric (in terms of F_1 -scores) is the Kulinski distance with an optimal k of 32. The F1 score of the optimal classifier is 0.774.

```
In [12]: # TODO: Fine-tune your model and report the best F1 score, use GridSearch
from sklearn.grid_search import GridSearchCV
metrics = ['jaccard', 'matching', 'dice', 'kulsinski', 'rogerstanimoto', 'russellrao']
k_list = map( lambda p: 2**p , range(6))
parameters = { 'n_neighbors':k_list, 'weights':['uniform', 'distance'], 'metric':metrics}
model = KNeighborsClassifier()
from sklearn.metrics import make_scorer
scorer = make_scorer( f1_score )
clf = GridSearchCV( model , parameters , scoring = "f1" )
clf.fit( X_train , y_train )
print clf.best_estimator_
print 'f1 score = {:.f}'.format(f1_score( clf.predict(X_test) , y_test))

KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='kulsinski',
                     metric_params=None, n_neighbors=32, p=2, weights='distance')
f1 score = 0.774194
```